

Assignment 4

Python Programming

Assignment Date	25 Oct 2022
Student Name	Naveenkumar S
Student Roll Number	923819104028
Maximum Marks	2 Marks

Question 1:

1 . Importing Required Package

Solution :

```
import pandas as pd
import numpy as np
import seaborn as sbn
import matplotlib.pyplot as plt
```

Question 2:

2 . Loading the Dataset

Solution :

```
db = pd.read_csv('/Mall_Customers.csv')
Db
```

Output

```
Out[4]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40
...
195	196	Female	35	120	79
196	197	Female	45	126	28
197	198	Male	32	126	74
198	199	Male	32	137	18
199	200	Male	30	137	83

: 200 rows Ã— 5 columns

Question 3:

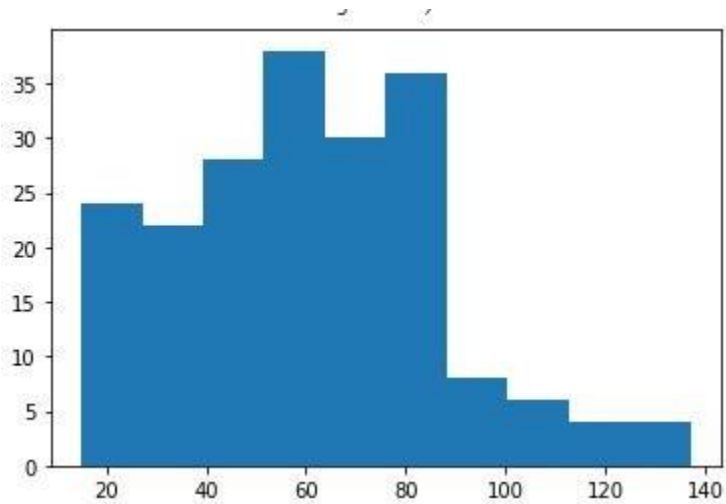
3 . Visualizations

3.1 UniVariate Analysis

3.1.1 Solution :

```
plt.hist(db['Annual Income (k$)'])
```

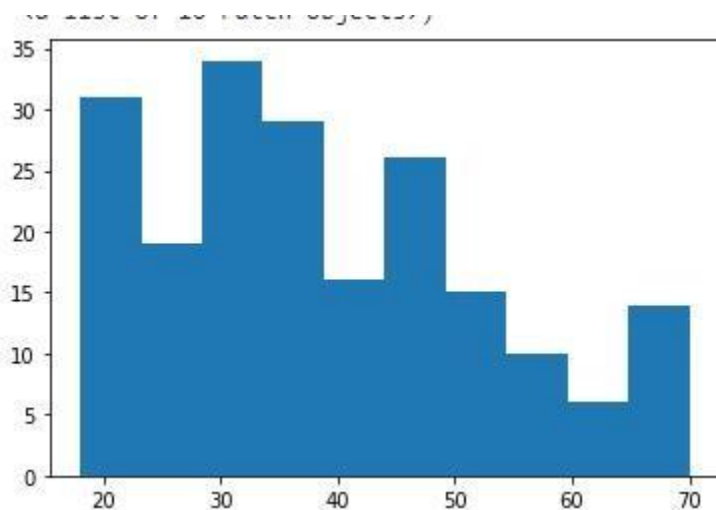
Output :



3.1.2 Solution

```
plt.hist(db['Age'])
```

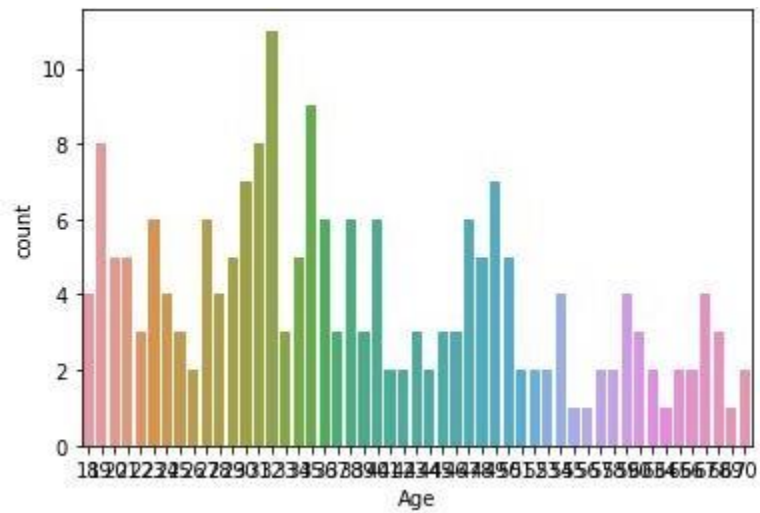
Output :



3.1.3 Solution :

```
sbn.countplot(db['Age'])
```

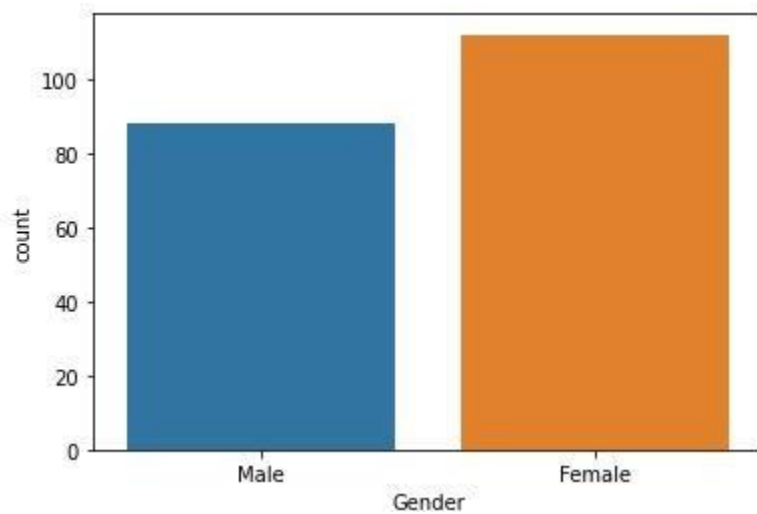
Output :



3.1.4 Solution :

```
sbn.countplot(db['Gender'])
```

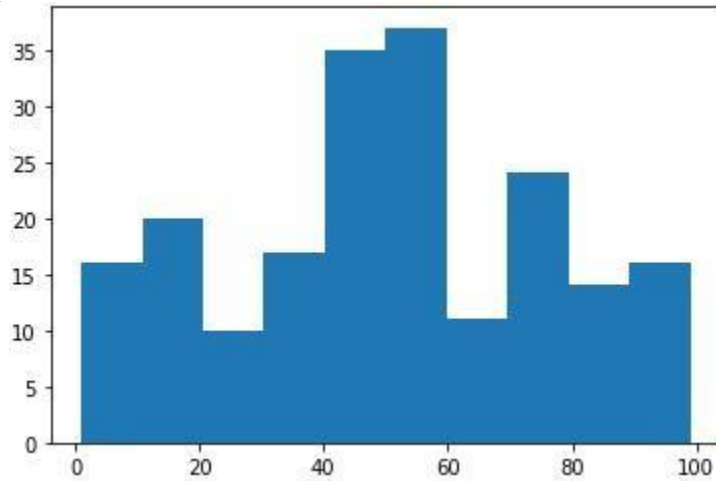
Output :



3.1.5 Solution :

```
plt.hist(db['Spending Score (1-100)'])
```

Output :

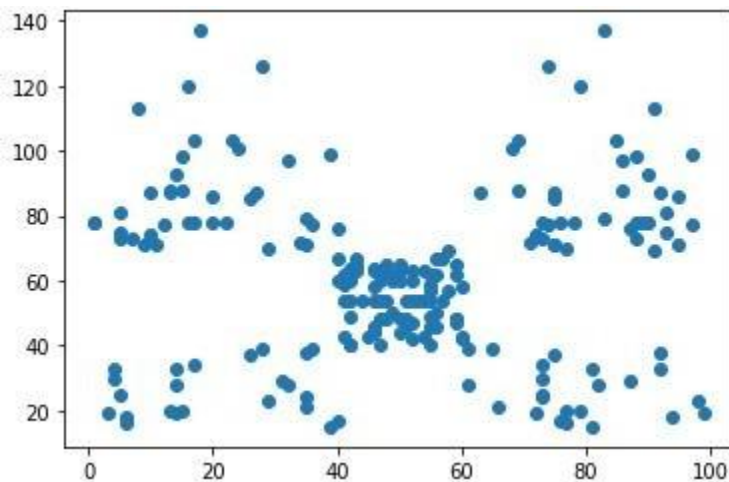


3.2 Bi-Variate Analysis

3.2.1 Solution :

```
plt.scatter(db['Spending Score (1-100)'],db['Annual Income (k$)'])
```

Output :



3.2.2 Solution :

```
plt.scatter(db['Gender'],db['Annual Income (k$)'])
```

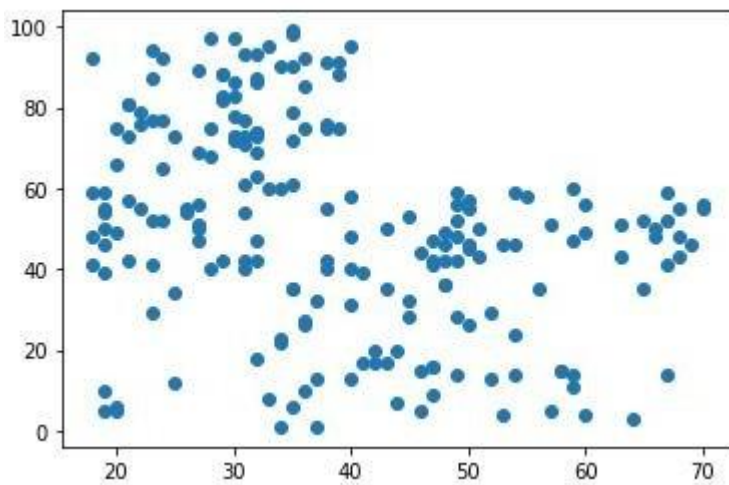
Output :



3.2.3 Solution :

```
plt.scatter(db['Age'],db['Spending Score (1-100)'])
```

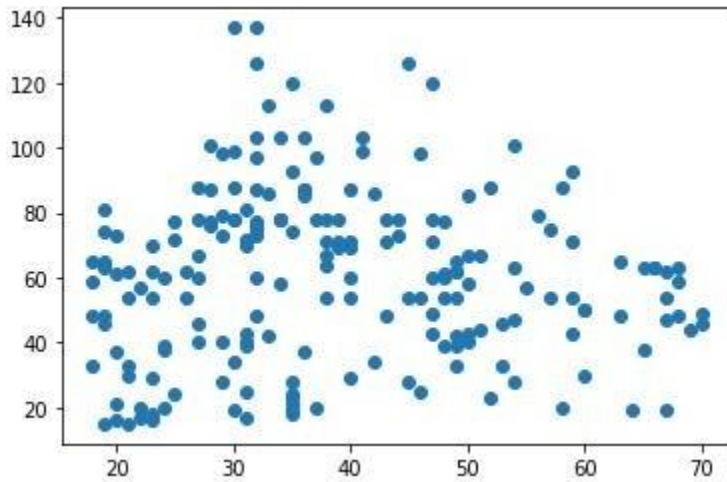
Output :



3.2.4 Solution :

```
plt.scatter(db['Age'],db['Annual Income (k$)'])
```

Output :



3.2.5 Solution :

```
sbn.heatmap(db.corr(), annot = True)
```

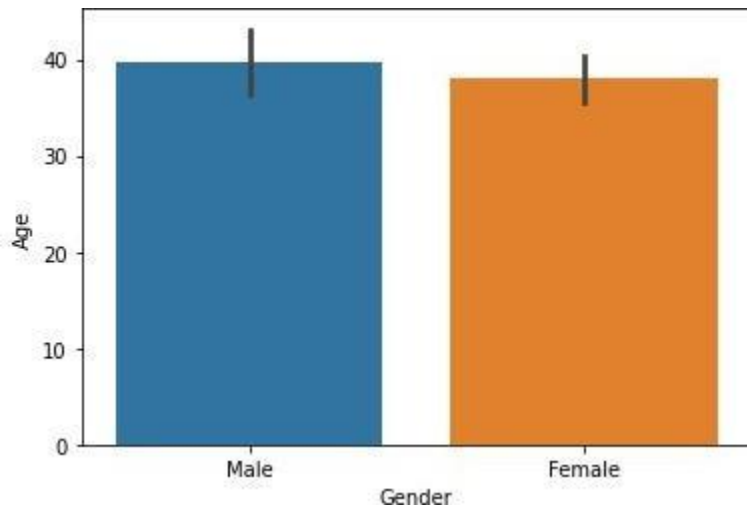
Output :



3.2.6 Solution :

```
sbn.barplot(db['Gender'], db['Age'])
```

Output :

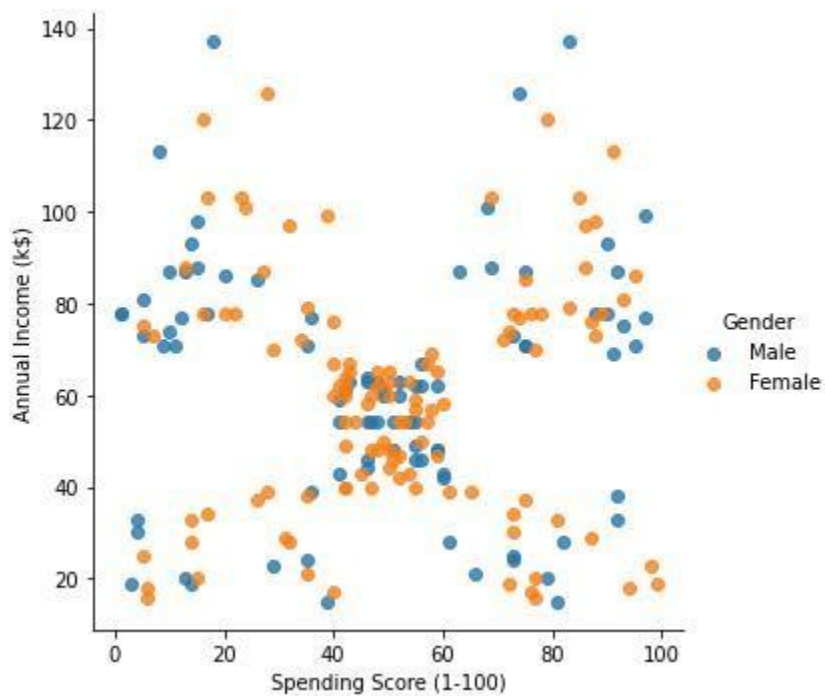


3.3 Multi-Variate Analysis

3.3.1 Solution :

```
sbn.lmplot("Spending Score (1-100)","Annual Income (k$)", db, hue="Gender", fit_reg=False);
```

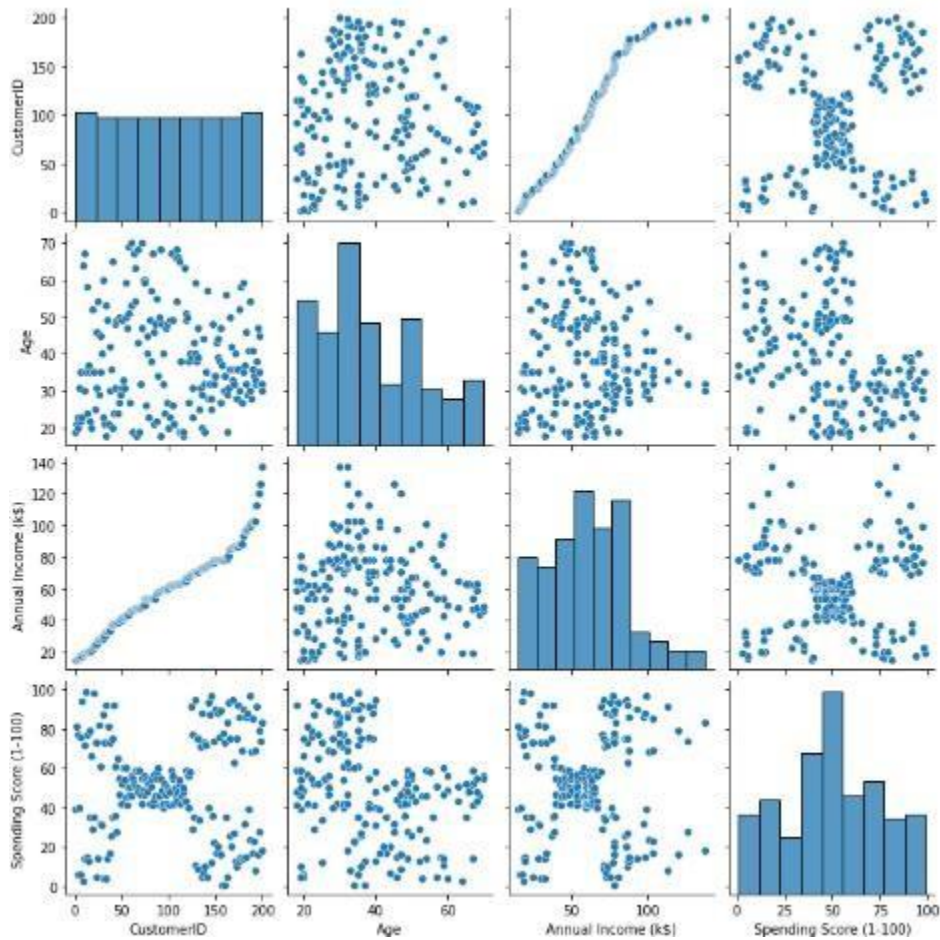
Output :



3.3.2 Solution :

```
sbn.pairplot(db)
```

Output :



Question 4:

4 . Perform descriptive statistics on the dataset

4.1 Solution :

```
db.describe()
```


Output :

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

4.2 Solution :

db.dtypes

Output :

```
CustomerID      int64
Gender          object
Age             int64
Annual Income (k$)  int64
Spending Score (1-100)  int64
dtype: object
```

4.3 Solution :

db.var()

Output :

```
] CustomerID      3350.000000
   Age           195.133166
   Annual Income (k$)  689.835578
   Spending Score (1-100)  666.854271
   dtype: float64
```

4.4 Solution :

db.skew()

Output :

```
CustomerID      0.000000
Age             0.485569
Annual Income (k$)  0.321843
Spending Score (1-100) -0.047220
dtype: float64
```

4.5 Solution :

db.corr()

Output :

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
CustomerID	1.000000	-0.026763	0.977548	0.013835
Age	-0.026763	1.000000	-0.012398	-0.327227
Annual Income (k\$)	0.977548	-0.012398	1.000000	0.009903
Spending Score (1-100)	0.013835	-0.327227	0.009903	1.000000

4.6 Solution :

db.std()

Output :

```
CustomerID      57.879185
Age             13.969007
Annual Income (k$)  26.264721
Spending Score (1-100)  25.823522
dtype: float64
```

Question 5:

5. Check for Missing values and deal with them

5.1 Solution :

db.isna().sum()

Output :

```
CustomerID      0
Gender          0
Age            0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64
```

5.2 Solution :

```
db.isna().sum().sum()
```

Output :

```
0
```

5.3 Solution :

```
db.duplicated().sum()
```

Output :

```
0
```

Question 6:

6 . Find the outliers and replace them outliers

6.1 Solution :

```
fig,ax=plt.subplots(figsize=(25,5))
```

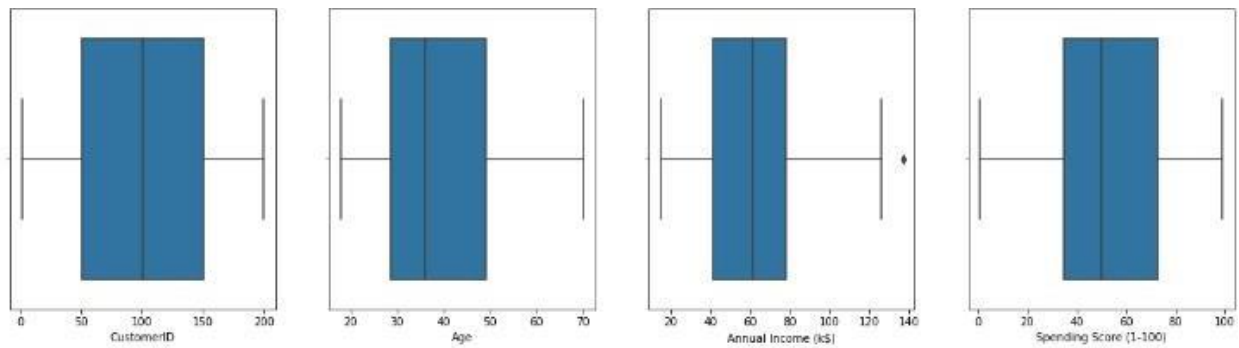
```
plt.subplot(1, 5, 2)
sbn.boxplot(x=db['Age'])
```

```
plt.subplot(1, 5, 3)
sbn.boxplot(x=db['Annual Income (k$)'])
```

```
plt.subplot(1, 5, 4)
sbn.boxplot(x=db['Spending Score (1-100)'])
```

```
plt.subplot(1, 5, 1)
sbn.boxplot(x=db['CustomerID'])
```

Output :



6.2 Solution :

```
quantile = db.quantile(q = [0.25, 0.75])
quantile
```

Output :

	CustomerID	Age	Annual Income (k\$)	Spending Score (1-100)
0.25	50.75	28.75	41.5	34.75
0.75	150.25	49.00	78.0	73.00

6.3 Solution :

```
quantile.loc[0.75]
```

Output :

```
CustomerID      150.25
Age              49.00
Annual Income (k$)  78.00
Spending Score (1-100)  73.00
Name: 0.75, dtype: float64
```

6.4 Solution :

```
quantile.loc[0.25]
```

Output :

```
CustomerID      50.75
Age             28.75
Annual Income (k$)  41.50
Spending Score (1-100)  34.75
Name: 0.25, dtype: float64
```

6.5 Solution :

```
IQR = quantile.iloc[1] - quantile.iloc[0]
IQR
```

Output :

```
CustomerID      99.50
Age             20.25
Annual Income (k$)  36.50
Spending Score (1-100)  38.25
dtype: float64
```

6.6 Solution :

```
upper = quantile.iloc[1] + (1.5 * IQR)
upper
```

Output :

```
CustomerID      299.500
Age             79.375
Annual Income (k$)  132.750
Spending Score (1-100)  130.375
dtype: float64
```

6.7 Solution :

```
lower = quantile.iloc[0] - (1.5 * IQR)
lower
```

Output :

```
CustomerID      -98.500
Age             -1.625
Annual Income (k$)  -13.250
Spending Score (1-100)  -22.625
dtype: float64
```

6.8 Solution :

```
db.mean()
```

Output :

```
CustomerID      100.50
Age              38.85
Annual Income (k$)  60.56
Spending Score (1-100)  50.20
dtype: float64
```

6.9 Solution :

```
db['Annual Income (k$)'].max()
```

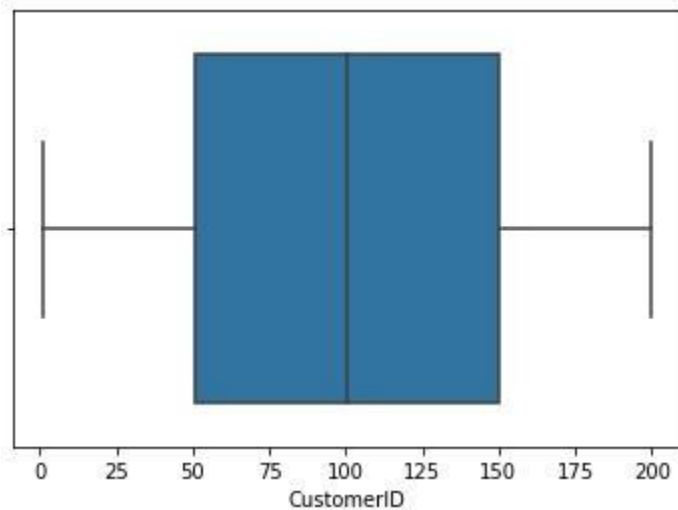
Output :

```
137
```

6.10 Solution :

```
sbn.boxplot(db['CustomerID'])
```

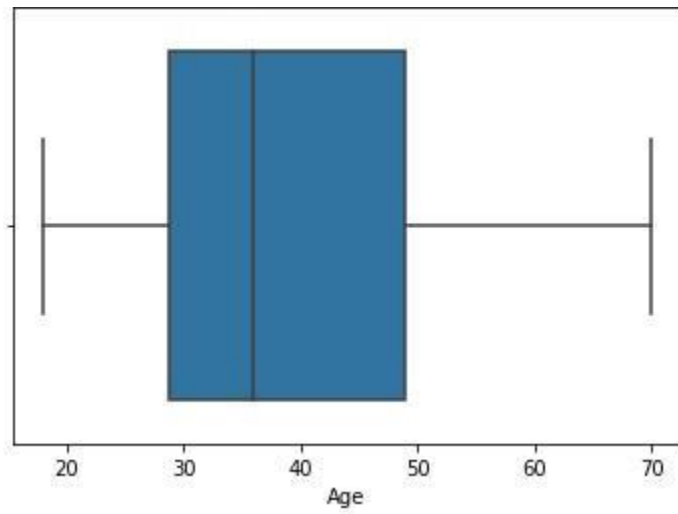
Output :



6.11 Solution :

```
sbn.boxplot(db['Age'])
```

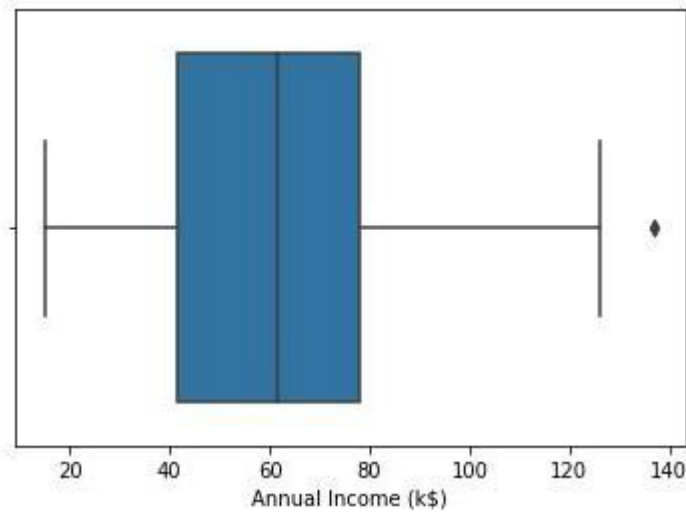
Output :



6.12 Solution :

```
sbn.boxplot(db['Annual Income (k$)'])
```

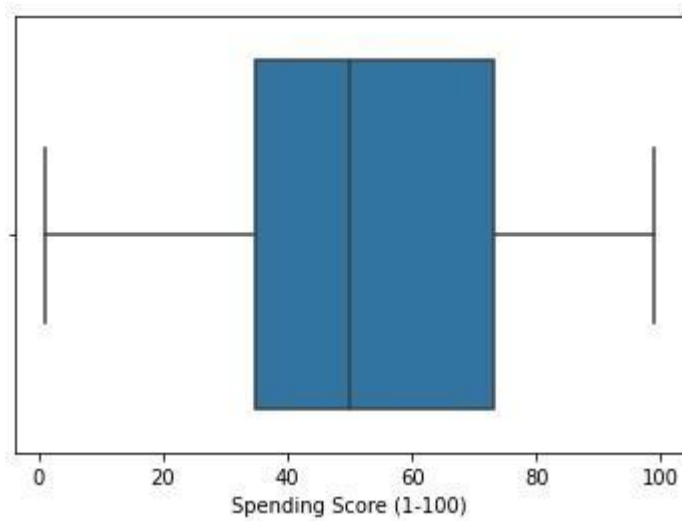
Output :



6.13 Solution :

```
sbn.boxplot(db['Spending Score (1-100)'])
```

Output :



Question 7:

7 . Check for Categorical columns and perform encoding

7.1 Solution :

```
db.select_dtypes(include='object').columns
```

Output :

```
Index(['Gender'], dtype='object')
```

7.2 Solution :

```
db['Gender'].unique()
```

Output :

```
array(['Male', 'Female'], dtype=object)
```

7.3 Solution :

```
db['Gender'].replace({'Male':1,'Female':0},inplace=True)  
db
```


Output :

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15.00	39
1	2	1	21	15.00	81
2	3	0	20	16.00	6
3	4	0	23	16.00	77
4	5	0	31	17.00	40
...
195	196	0	35	120.00	79
196	197	0	45	126.00	28
197	198	1	32	126.00	74
198	199	1	32	60.55	18
199	200	1	30	60.55	83

200 rows Ã— 5 columns

7.4 Solution :

db.head()

Output :

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15.0	39
1	2	1	21	15.0	81
2	3	0	20	16.0	6
3	4	0	23	16.0	77
4	5	0	31	17.0	40

Question 8:

8 . Scaling the data

8.1 Solution :

```
from sklearn.preprocessing import StandardScaler
ss = StandardScaler().fit_transform(db)
```

Output :

Output :

Question 9:

9. Perform any of the clustering algorithms

9.1 Solution :

```
from sklearn.cluster import KMeans
TWSS = []
k = list(range(2,9))
```

```

for i in k:
    kmeans = KMeans(n_clusters = i , init = 'k-means++')
    kmeans.fit(db)
    TWSS.append(kmeans.inertia_)
TWSS

```

Output :

```

[381507.64738523855,
 268062.55433747417,
 191557.78099047023,
 153327.3825004856,
 119166.15727643928,
 101296.86197582977,
 85792.73210128325]

```

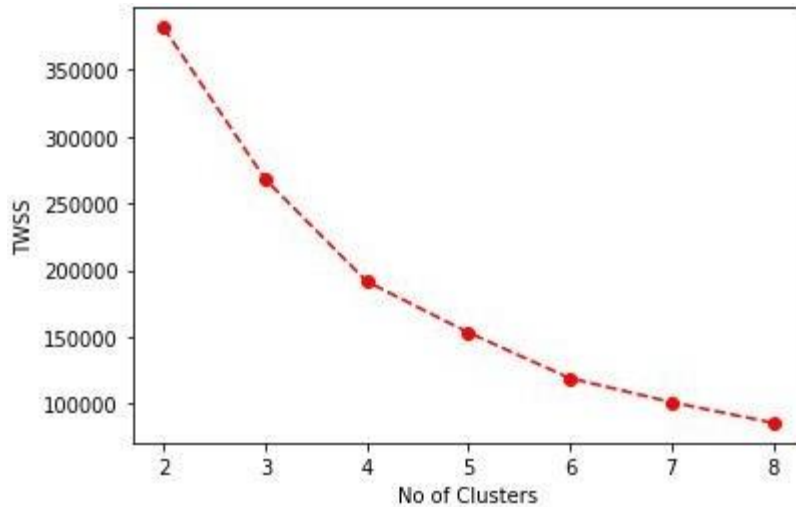
9.2 Solution :

```

plt.plot(k,TWSS, 'ro--')
plt.xlabel('No of Clusters')
plt.ylabel('TWSS')

```

Output :



9.3 Solution :

```

model = KMeans(n_clusters = 4)
model.fit(db)

```

Output :

```

KMeans(n_clusters=4)

```

9.4 Solution :

```
mb = pd.Series(model.labels_)
db['Cluster'] = mb
db
```

Output :

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	Cluster
0	1	1	19	15.00	39	1
1	2	1	21	15.00	81	1
2	3	0	20	16.00	6	1
3	4	0	23	16.00	77	1
4	5	0	31	17.00	40	1
...
195	196	0	35	120.00	79	2
196	197	0	45	126.00	28	0
197	198	1	32	126.00	74	2
198	199	1	32	60.55	18	0
199	200	1	30	60.55	83	2

200 rows Ã— 6 columns

9.5 Solution :

```
mb=pd.Series(model.labels_)
db.head(3)
```

Output :

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	Cluster
0	1	1	19	15.0	39	1
1	2	1	21	15.0	81	1
2	3	0	20	16.0	6	1

Question 10:

10 . Add the cluster data with the primary dataset

10.1 Solution :

```
db['Cluster']=kmeans.labels_  
db.head()
```

Output :

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	Cluster
0	1	1	19	15.0	39	5
1	2	1	21	15.0	81	2
2	3	0	20	16.0	6	5
3	4	0	23	16.0	77	2
4	5	0	31	17.0	40	5

10.2 Solution :

```
db.tail()
```

Output :

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	Cluster
195	196	0	35	120.00	79	6
196	197	0	45	126.00	28	1
197	198	1	32	126.00	74	6
198	199	1	32	60.55	18	1
199	200	1	30	60.55	83	6

Question 11:

11 . Split the data into dependent and independent variables

11.1 Solution :

```
X=db.drop('Cluster',axis=1)  
Y=db['Cluster']  
y=db['Cluster']  
y
```

Output :

```
0      5
1      2
2      5
3      2
4      5
..
195    6
196    1
197    6
198    1
199    6
Name: Cluster, Length: 200, dtype: int32
```

11.2 Solution :

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,Y,test_size=0.2,random_state=42)

print("Number transactions X_train dataset: ", X_train.shape)
print("Number transactions y_train dataset: ", y_train.shape)
print("Number transactions X_test dataset: ", X_test.shape)
print("Number transactions y_test dataset: ", y_test.shape)
```

Output :

```
Number transactions X_train dataset: (160, 5)
Number transactions y_train dataset: (160,)
Number transactions X_test dataset: (40, 5)
Number transactions y_test dataset: (40,)
```

Question 12:

12 . Split the data into training and testing

12.1 Solution :

X_train

Output :

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
79	80	0	49	54.0	42
197	198	1	32	126.0	74
38	39	0	36	37.0	26
24	25	0	54	28.0	14
122	123	0	40	69.0	58
...
106	107	0	66	63.0	50
14	15	1	37	20.0	13
92	93	1	48	60.0	49
179	180	1	35	93.0	90
102	103	1	67	62.0	59

160 rows Ã— 5 columns

12.2 Solution :

X_test

Output :

CustomerID	Gender	Age	Annual Income (k\$)	SpendingScore (1-100)	
95	96	1	24	610	53
19	70	1	22	202	73
30	21	1	60	103	4
58	150	1	34	785	1
28	120	1	59	214	11
15	110	0	19	650	50
09	70	0	32	480	47
70	171	1	40	670	11
74	175	0	52	880	13
45	40	0	24	390	65
00	07	0	43	480	50
82	183	1	46	980	75
65	106	0	36	850	75
78	79	0	25	540	52
86	187	0	54	1010	24
77	178	1	27	880	60
56	57	0	57	440	90
52	153	0	44	780	20
82	83	1	67	540	41
08	60	1	19	480	58
24	123	0	23	700	23
16	17	0	35	270	35
48	140	0	34	780	22
93	94	0	40	840	46
65	66	1	18	480	95
60	61	1	70	460	56
04	05	0	27	540	53
67	68	0	68	480	48
25	126	0	31	700	77
32	133	0	25	720	34
9	10	0	30	190	72
18	19	1	52	230	25
55	58	1	47	430	47
75	76	1	26	540	54
50	151	1	43	780	77
04	105	1	49	820	58
35	136	0	29	730	88
37	138	1	32	730	73
164	165	1	50	850	26
76	77	0	45	540	55

12.3 Solution :

y_train

Output :

```
79      4
197     6
38      5
24      5
122     0
..
106     0
14      5
92      0
179     6
102     0
Name: Cluster, Length: 160, dtype: int32
```

12.14 Solution :

y_test

Output :

95	0
15	2
30	5
158	7
128	7
115	0
69	4
170	1
174	1
45	2
66	4
182	1
165	6
78	0
186	1
177	6
56	4
152	7
82	4
68	4
124	7
16	5
148	7
93	0
65	4
60	4
84	0
67	4
125	3
132	7
9	2
18	5
55	4
75	4
150	7
104	0
135	3
137	3
164	1
76	4

Name: Cluster, dtype: int32

Question 13:

13 . Build the Model

13.1 Solution :

```

from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
model.fit(X_train,y_train)
from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
model.fit(X_train,y_train)

```

Output :

```
LogisticRegression()
```

Question 14:

14 . Train the Model

Solution :

```
model.score(X_train,y_train)
```

Output :

```
0.83125
```

Question 15:

15 . Test the Model

Solution :

```
model.score(X_test,y_test)
```

Output :

```
0.675
```

Question 16:

16 . Measure the performance using Evaluation Metrics

16.1 Solution :

```
from sklearn.metrics import confusion_matrix,classification_report
y_pred=model.predict(X_test)
confusion_matrix(y_test,y_pred)
```

Output :

```
array([[5, 0, 0, 0, 0, 0, 1, 0],
       [0, 5, 0, 0, 0, 0, 0, 0],
       [0, 0, 3, 0, 0, 0, 0, 0],
       [0, 0, 0, 3, 0, 0, 0, 0],
       [3, 0, 2, 0, 6, 0, 0, 0],
       [0, 0, 0, 0, 0, 3, 0, 0],
       [0, 0, 0, 1, 0, 0, 1, 0],
       [0, 6, 0, 0, 0, 0, 0, 1]])
```

16.2 Solution :

```
print(classification_report(y_test,y_pred))
```

Output :

	precision	recall	f1-score	support
0	0.62	0.83	0.71	6
1	0.45	1.00	0.62	5
2	0.60	1.00	0.75	3
3	0.75	1.00	0.86	3
4	1.00	0.55	0.71	11
5	1.00	1.00	1.00	3
6	0.50	0.50	0.50	2
7	1.00	0.14	0.25	7
accuracy			0.68	40
macro avg	0.74	0.75	0.68	40
weighted avg	0.80	0.68	0.64	40